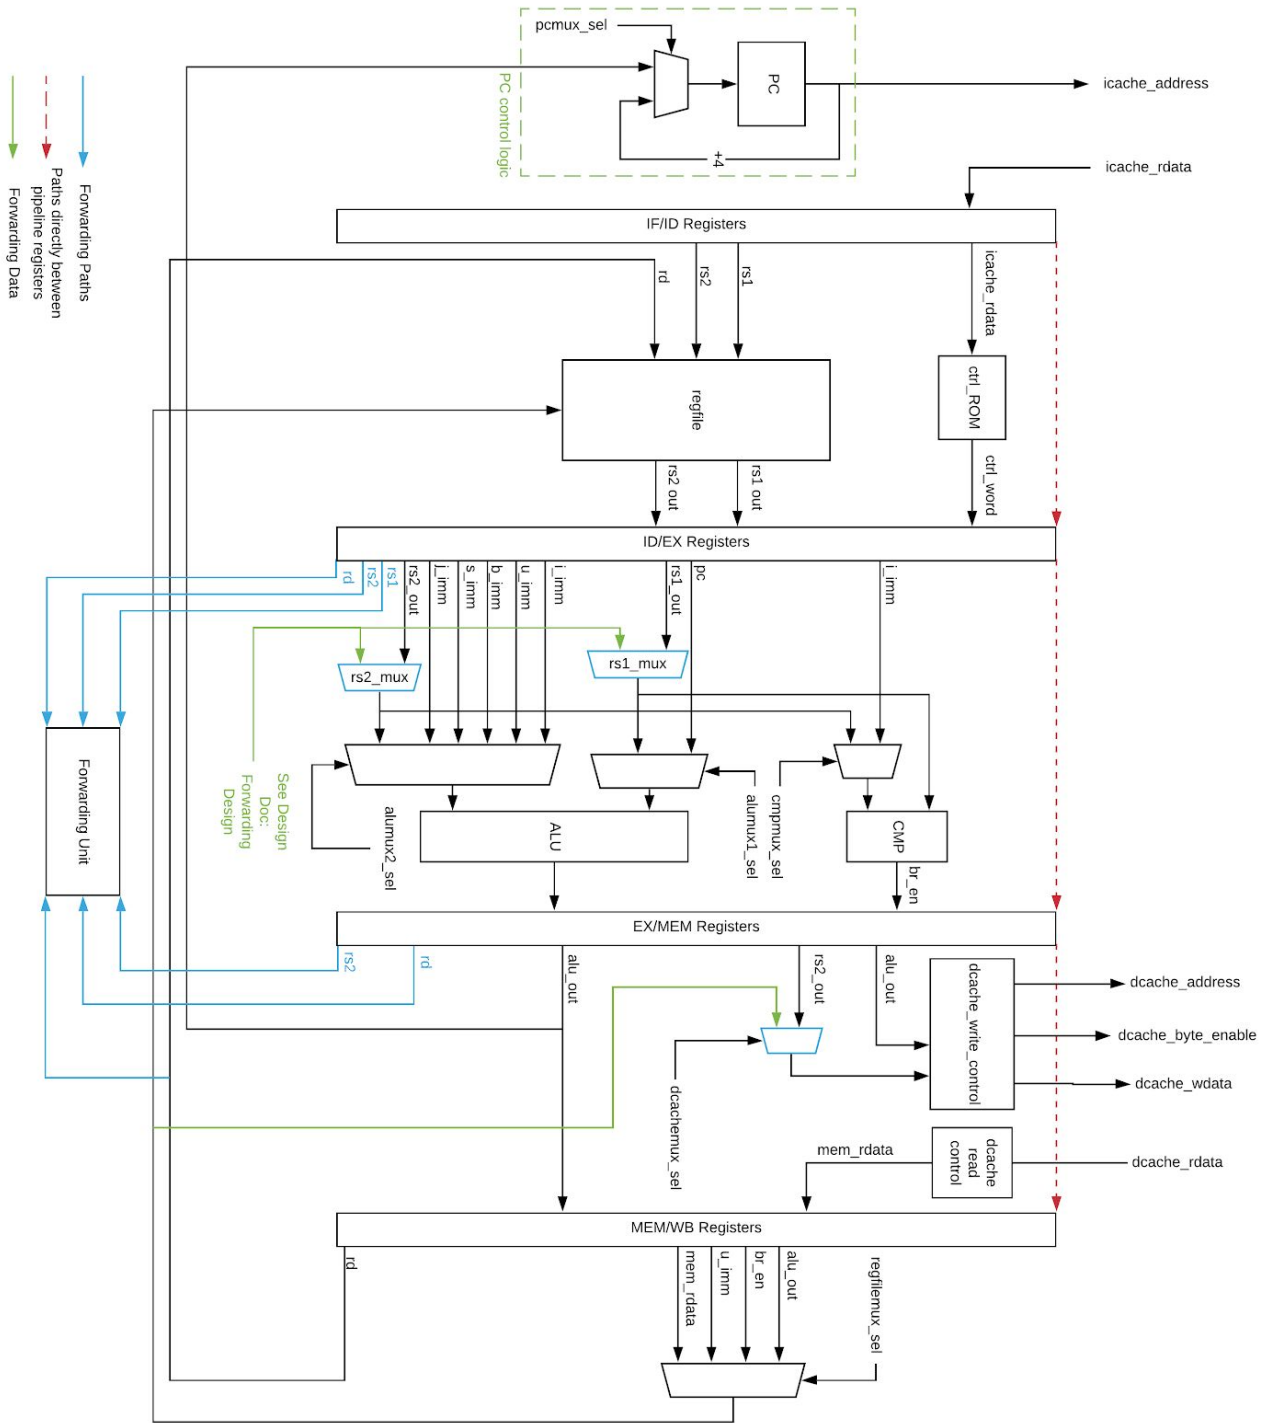


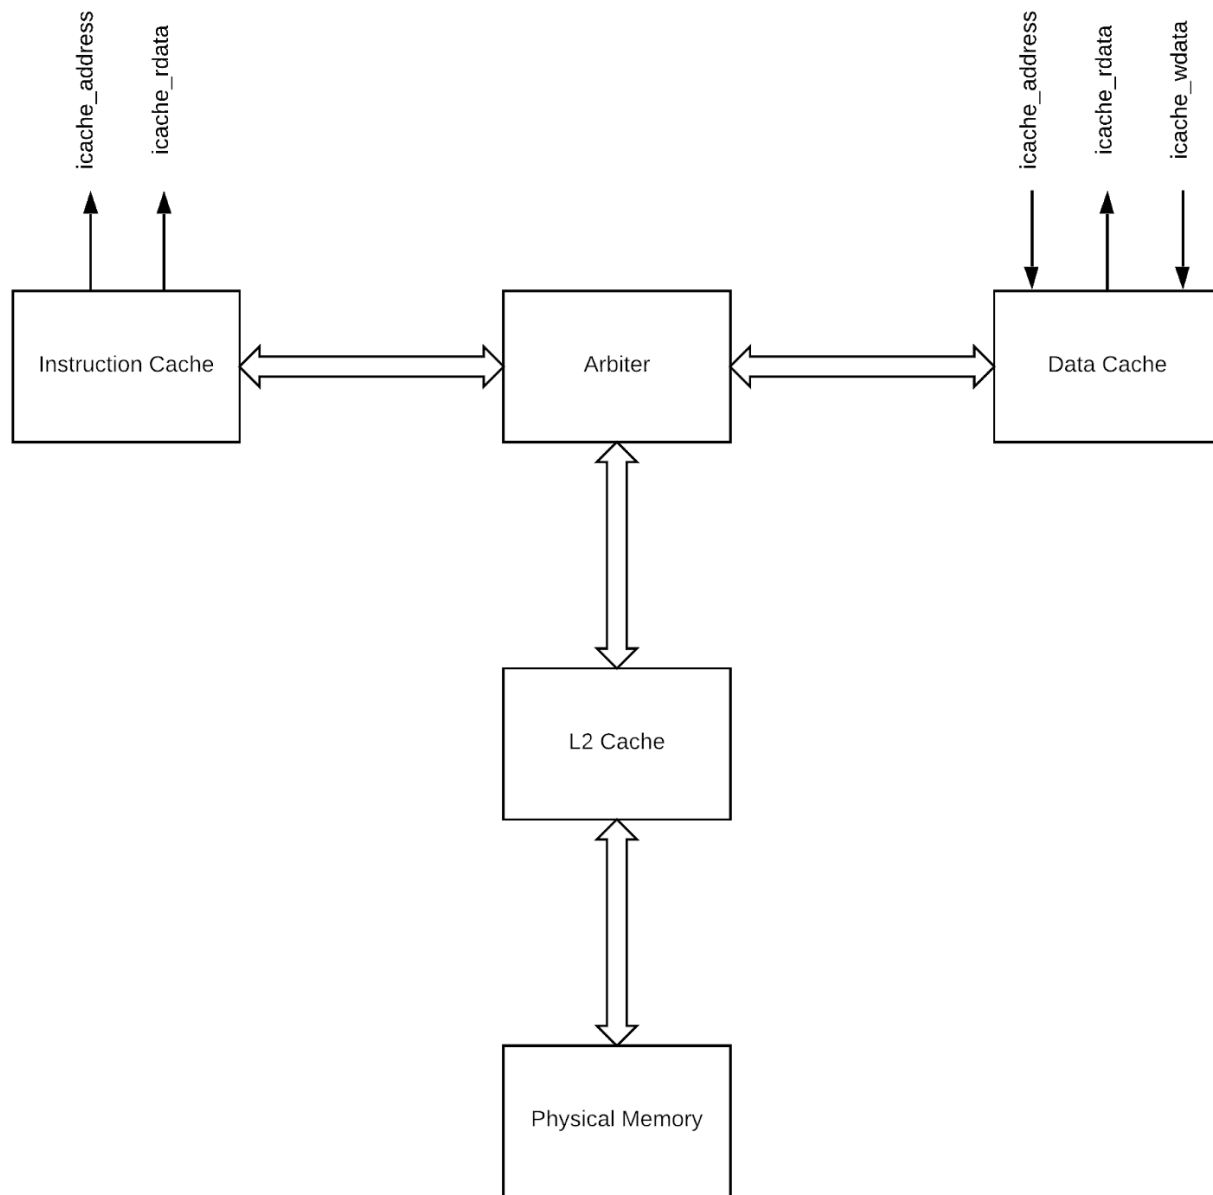
# ECE 411 MP3 Design Specification:

## Group ZBA

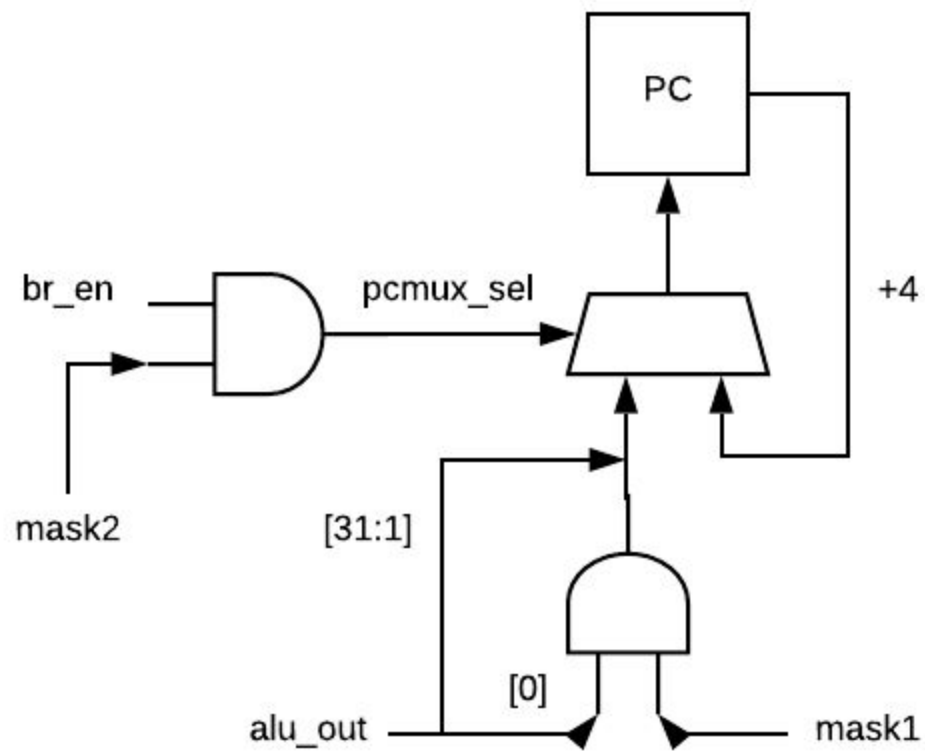
Zerui “Jerry” An, Austin Smithe, Brock Boehler



# Memory Hierarchy



## PC Control Logic



$\text{mask1} = \sim(\text{opcode} == \text{jlr})$   
 $\text{mask2} = (\text{opcode} == \text{br/jal/jalr})$

## Pipeline Registers

### IF/ID

Register name	Description
icache_rdata	32-bit instruction
PC	Program counter

### ID/EX

Register name	Description
icache_rdata	32-bit instruction
ctrl_word	A struct containing all control signals. Output of control ROM.
PC	Program counter

### EX/MEM

Register Name	Description
alu_out	ALU output
rs2_out	Contents of register 2
rd	Destination register
br_en	Branch enable
u_imm	Unsigned immediate from instruction data
ctrl_word	A struct containing all control signals.

### MEM/WB

Register Name	Description
alu_out	ALU output
rd	Destination register

br_en	Branch enable
u_imm	Unsigned immediate from instruction data
MDR_out	Data received from data cache
ctrl_word	A struct containing all control signals.i

## Control Signals

### IF

Signal Name	Description
load_pc	Load program counter
pcmux_sel	Selects between PC+4 or WB stage output

### ID

Signal Name	Description
<i>None</i>	

### EX

Signal Name	Description
rs1mux_sel	Output of forwarding unit. Used to select the signal to be forwarded.
rs2mux_sel	
alumux1_sel	ALU input 1 select
alumux2_sel	ALU input 2 select
cmpmux_sel	Chooses input to CMP module
aluop	ALU operation
cmpop	Compare operation

## MEM

Signal Name	Description
dcachemux_sel	Selects input to write to data cache

## WB

Signal Name	Description
regfilemux_sel	Chooses data to be written to regfile
load_regfile	Loads regfile

## Information Shared Between Pipeline Registers

IF/ID → ID/EX	ID/EX → EX/MEM	EX/MEM → MEM/WB
icache_rdata PC ctrl_word	rd u_imm ctrl_word	alu_out rd br_en u_imm ctrl_word

ctrl\_word struct:

Signal Name	Length	Stage
alumux1_sel	1	EX
alumux2_sel	3	
cmpmux_sel	1	
aluop	3	
cmpop	3	
dcache_read	1	MEM

dcache_write	1	
mask1	1	
mask2	1	
regfilemux_sel	3+	WB
load_regfile	1	

## Control Word Logic:

*Note: STB = See Table Below*

Signal Name	lui	auipc	jal	jalr	br	load	store	imm/reg
alumux1_sel	x	pc	pc	rs1_out	pc	rs1_out	rs1_out	rs1_out
alumux2_sel	x	u_imm	j_imm	i_imm	b_imm	i_imm	s_imm	STB
cmpmux_sel	x	x	x	x	rs2_out	x	x	STB
aluop	x	alu_add	alu_add	alu_add	alu_add	alu_add	alu_add	STB
cmpop	x	x	x	x	funct3	x	x	STB
dcache_read	0	0	0	0	0	1	0	0
dcache_write	0	0	0	0	0	0	1	0
regfilemux_sel	u_imm	alu_out	pc_plus4	pc_plus4	x	MDR_out	x	STB
load_regfile	1	1	1	1	0	1	0	STB

Signal Name	imm(default)	imm(srai)	imm(slti)	imm(sltiu)
alumux1_sel	rs1_out			
alumux2_sel	i_imm	i_imm	x	x
cmpmux_sel	x	x	i_imm	i_imm
aluop	funct3	alu_sra	x	x



cmpop	x	x	blt	bltu
dcache_read	0			
dcache_write	0			
regfilemux_sel	alu_out	alu_out	br_en	br_en
load_regfile	1			

Signal Name	reg(default)	reg(sra)	reg(slt)	reg(sltu)	reg(sub)
alumux1_sel	rs1_out				
alumux2_sel	rs2_out	rs2_out	x	x	rs2_out
cmpmux_sel	x	x	rs2_out	rs2_out	x
aluop	funct3	alu_sra	x	x	alu_sub
cmpop	x	x	blt	bltu	x
dcache_read	0				
dcache_write	0				
regfilemux_sel	alu_out	alu_out	br_en	br_en	alu_out
load_regfile	1				

## Read / Write Logic

Instruction	dcache_address & 0x03	dcache_mem_byte_enable
SW	x	4'b1111
SH	0	4'b0011
	1	4'b0110
	2	4'b1100
	3	x

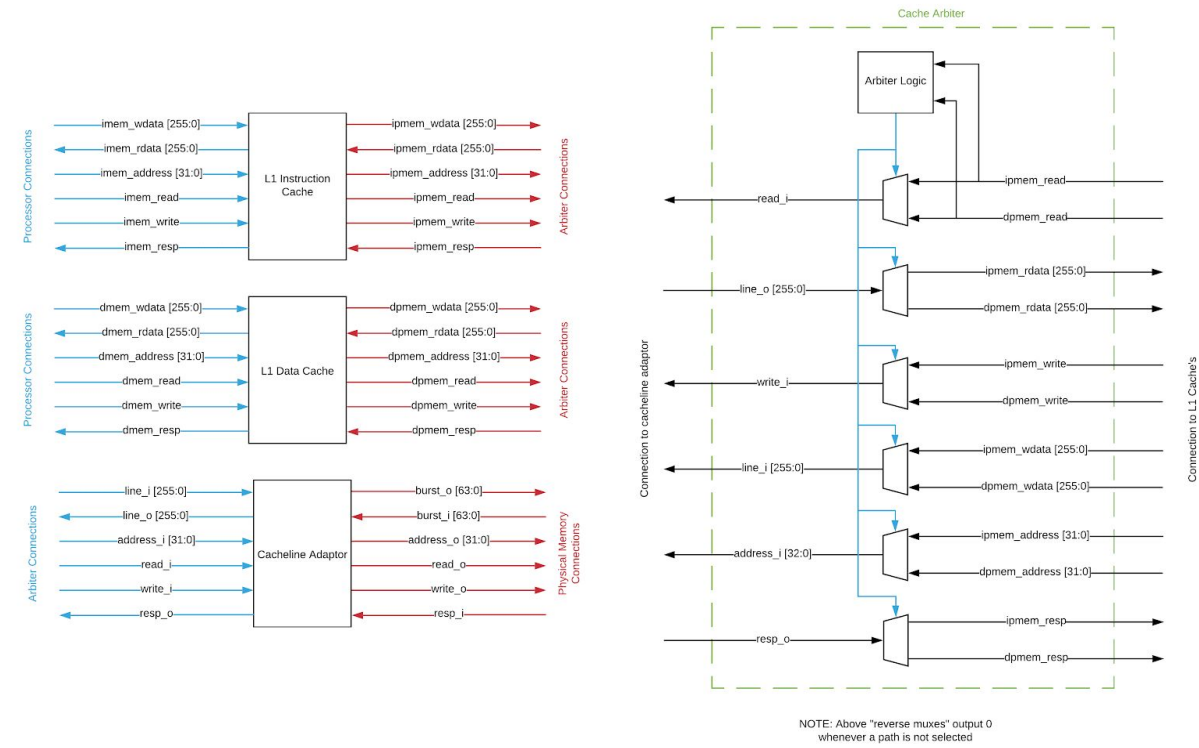
SB	0	4'b0001
	1	4'b0010
	2	4'b0100
	3	4'b1000

Instruction	dcache_wdata
SW	dcache_wdata <= rs2_out
SH	dcache_wdata <= (rs2_out << {alu_out[1:0]}, 3'd0))
SB	dcache_wdata <= (rs2_out << {alu_out[1:0]}, 3'd0))


instruction	dcache_address & 0x03	MDRreg_out
lw	x	dcache_rdata
lh	0	{{16{dcache_rdata[15]}}, dcache_rdata[15:0]}
	1	{{16{dcache_rdata[23]}}, dcache_rdata[23:8]}
	2	{{16{dcache_rdata[31]}}, dcache_rdata[31:16]}
	3	x
lhu	0	{16'd0, dcache_rdata[15:0]}
	1	{16'd0, dcache_rdata[23:8]}
	2	{16'd0, dcache_rdata[31:16]}
	3	x
lb	0	{{24{dcache_rdata[7]}}, dcache_rdata[7:0]}
	1	{{24{dcache_rdata[15]}}, dcache_rdata[15:8]}
	2	{{24{dcache_rdata[23]}}, dcache_rdata[23:16]}
	3	{{24{dcache_rdata[31]}}, dcache_rdata[31:24]}

Ibu	0	{24'd0, dcache_rdata[7:0]}
	1	{24'd0, dcache_rdata[15:8]}
	2	{24'd0, dcache_rdata[23:16]}
	3	{24'd0, dcache_rdata[31:24]}

# Cache Arbiter



## Cache Arbiter Control / State Machine



Note: The cache arbiter prioritizes the instruction cache over the data cache. Whenever both want to read at the same time it is the instruction cache that gets to read / write.

Transition	Condition
A	<code>(dpmem_read == 1'b0 &amp; dpmem_write == 1'b0) &amp; (ipmem_read == 1'b1   ipmem_write == 1'b1) &amp; (resp_o == 1'b1)</code>
B	<code>(ipmem_read == 1'b0 &amp; ipmem_write == 1'b0) &amp; (dpmem_read == 1'b1   dpmem_write == 1'b1) &amp; (resp_o == 1'b1)</code>
C	<code>(dpmem_read == 1'b1   dpmem_write == 1'b1) &amp; (ipmem_read == 1'b0 &amp; ipmem_write == 1'b0)</code>
D	<code>(ipmem_read == 1'b1   ipmem_write == 1'b1)</code>
E	<code>(dpmem_read == 1'b0 &amp; dpmem_write == 1'b0) &amp; (ipmem_read == 1'b0 &amp; ipmem_write == 1'b0)</code>
F	<code>(dpmem_read == 1'b0 &amp; dpmem_write == 1'b0)</code>

	<code>(ipmem_read == 1'b0 &amp; ipmem_write == 1'b0) &amp;</code>
<b>G</b>	<code>(dpmem_read == 1'b0 &amp; dpmem_write == 1'b0) &amp;</code> <code>(ipmem_read == 1'b0 &amp; ipmem_write == 1'b0)</code>
<b>H</b>	<code>(dpmem_read == 1'b1   dpmem_write == 1'b1) &amp;</code> <code>(ipmem_read == 1'b0 &amp; ipmem_write == 1'b0) &amp;</code> <code>(resp_o == 1'b1)</code>
<b>I</b>	<code>(ipmem_read == 1'b1   ipmem_write == 1'b1) &amp;</code> <code>(resp_o == 1'b1)</code>
<b>J</b>	<code>rst == 1'b1</code>

State	Output
idle	Don't Care
data	Muxes output data cache signals
instruction	Muxes output instruction cache signals.

## Forwarding Design

There are hazard between the following stages:

### EX/MEM ⇒ EX

```
if (
    exmem_load_regfile == 1
&&  exmem_opcode != load
&&  exmem_rd == idex_rs
)
```

Example:

```
ADD x1, x2, x2;
ADD x3, x1, x1;
```

### MEM/WB ⇒ EX

```
if (
    memwb_load_regfile == 1
&&  memwb_rd == idex_rs
)
```

**Example:**

```
ADD x1, x2, x2;
```

```
NOP
```

```
ADD x3, x1, x1;
```

```
if (
    exmem_load_regfile == 1
    && exmem_opcode == load
    && !(idx_opcode == store && exmem_rd != idx_rs1)
    && exmem_rd == idx_rs
)
```

**Example(2):**

```
LW x1, ADDR;
```

```
ADD x3, x1, x1; (1-cycle stall)
```

**MEM/WB ⇒ MEM****Example:**

```
LW x1, ADDR1;
```

```
SW x1, ADDR2;
```

The forwarding muxes, pictured in the datapath, have the following connections:

Forwarding Mux	Inputs
rs1_mux	rs1_out
	EXMEM_ALU_out
	EXMEM_BR_EN
	regfilemux_out
	mem_rdata
	EXMEM_u_imm
rs2_mux	rs2_out
	EXMEM_ALU_out
	EXMEM_BR_EN
	regfilemux_out
	mem_rdata

	EXMEM_u_imm
dcache_mux	rs2_out
	regfilemux_out

The hazards and their forwarding solutions are as follows:

#### **rs1\_idex == rd\_exmem**

Between Stages: EX/MEM  $\Rightarrow$  EX

<b>EX/MEM Instruction</b>	<b>rs1_mux</b>
LUI	EXMEM_u_imm
AUIPC	EXMEM_ALU_out
branch	x
load	mem_rdata
store	x
SLTI	EXMEM_BR_EN
SLTUI	EXMEM_BR_EN
IMM (other)	EXMEM_ALU_out
SLT	EXMEM_BR_EN
SLTU	EXMEM_BR_EN
REG (other)	EXMEM_ALU_out

#### **rs2\_idex == rd\_exmem**

Between Stages: EX/MEM  $\Rightarrow$  EX

<b>EX/MEM Instruction</b>	<b>rs2_mux</b>
LUI	EXMEM_u_imm
AUIPC	EXMEM_ALU_out
branch	x
load	mem_rdata

store	x
SLTI	EXMEM_BR_EN
SLTUI	EXMEM_BR_EN
IMM (other)	EXMEM_ALU_out
SLT	EXMEM_BR_EN
SLTU	EXMEM_BR_EN
REG (other)	EXMEM_ALU_out

#### rs1\_idex == rd\_memwb

Between Stages: MEM/WB  $\Rightarrow$  EX

MEM/WB Instruction	rs1_mux
<i>Any</i>	regfilemux_out

#### rs2\_idex == rd\_memwb

Between Stages: MEM/WB  $\Rightarrow$  EX

MEM/WB Instruction	rs2_mux
<i>Any</i>	regfilemux_out

#### rs2\_exmem == rd\_memwb

Between Stages: MEM/WB  $\Rightarrow$  MEM

MEM/WB Instruction	dcache_mux
<i>Any</i>	regfilemux_out

#### Important Notes:

1. Whenever there is a hazard with both the memory stage and writeback stage needing to write to the same register in the execute stage, the instruction in the memory stage takes precedence.
2. The mem\_rdata mentioned above is not the data directly from dcache, which would create a very long critical path if used. Instead, we load the data from dcache in a register first, and use the register output in forwarding after a 1-cycle stall.




## Advanced Design Options

8-way pLRU L2 cache:

Interface of pLRU\_tree:

Signal Name	Direction	Notes
clk	input	
rst	input	
load	input	Signals the module to update its internal states (L6:L0)
index	input	Index of the set to update
last_access	input	The most-recent accessed way
plru	output	The pseudo-LRU



Our pLRU design would be a tree-like structure, where each bit of L[6:0] is kept in a tree node. A tree node would be similar to a register, but in addition to the standard I/O of a register array (clk, rst, load, index, data\_in, data\_out), it also has 2 1-bit output: load\_left and load\_right, which are the load signals that goes to its children.

For simplicity, we'll ignore the "index" input in the following discussion, as the pLRU-tree of each set functions identically and independently.

Update Rule:

- (1) L0 (tree root) is always updated to last\_access[2](MSB)
- (2) If a node is updated (load = 1), one of its children also needs to be updated.
- (3) If a node is not updated (load = 0), none of its children will be updated.


Thus, load\_left and load\_right is calculated as follows:

- (1) load\_left = load & (~data\_in) // Go left if data\_in is 0
- (2) load\_right = load & data\_in // Go right if data\_in is 1

Data\_in of each node is assigned as follows:

- (1) last\_access[2] (MSB) goes to L0 (root level)
- (2) last\_access[1] goes to L1 and L2
- (3) last\_access[0] (LSB) goes to L3, L4, L5, and L6 (leaf level)

Pseudo-LRU is computed by the following circuit:



## Eviction Write Buffer

### Signals:

Ports to L2 cache: *read\_i*, *write\_i*, *addr\_i*, *rdata\_o*, *wdata\_i*, *resp\_o*

Ports to pmem (cacheline adapter): *read\_o*, *write\_o*, *addr\_o*, *rdata\_i*, *wdata\_i*, *resp\_i*

Internal data: *wdata*, *addr*, *valid*.

### State Actions and Transitions:

#### Idle:

If (*read\_i* == 1):

If requested data matches the cacheline stored in EWB, set *rdata\_o* = *wdata*, and assert *resp\_o*. Otherwise, directly connect L2 cache to pmem ports.

If (*write\_i* == 1):

Load *addr\_i* and *wdata\_i* into internal registers (*addr* and *wdata*), then assert *resp\_o*.

Go to **Wait** state.

#### Wait:

(Idea: because we are using write-back cache, and a memory write is performed just now, there's very likely a memory read immediately after. Write-back will start after the next read completes)

Directly connect L2 cache (read interface) to pmem ports.

Wait for one of the following:

- L2 finishes reading from memory
- Several cycles of inactivity
- L2 asserts *write\_i*

Then, go to **Write** state.

#### Write:

Write buffered data back to memory.

Can serve EWB-hit reads at this state.

Ignores all other r/w requests from L2 cache.

Go to **Idle** state once write is completed.

## RISC-V M Extension


### Integration into Datapath

Integration into the datapath is trivial - divider and multiplier circuits will be added into the execution stage of the pipeline and used much like the ALU, and the control signals will be modified to accommodate the new hardware.

Due to the slow, complex nature of multiplication, and especially division, may force the datapath to be halted whenever waiting for a mathematical result from those units.

### Multiplier

The multiplier is a modified “Wallace Tree” multiplier, which is based off of lattice multiplication:



*Image Source:* Javier Rosa, *Step 3 of the lattice (shabakh) multiplication algorithm*, 7 March 2014, SVG Image, Wikipedia, accessed 20 April 2020,

[https://en.wikipedia.org/wiki/Lattice\\_multiplication#/media/File:Example\\_of\\_step\\_3\\_of\\_lattice\\_\(shabakh\)\\_multiplication\\_algorithm.svg](https://en.wikipedia.org/wiki/Lattice_multiplication#/media/File:Example_of_step_3_of_lattice_(shabakh)_multiplication_algorithm.svg)

In this method of multiplication, the two large numbers are multiplied digit-wise in a table, the overflow and result are organized in the way shown above, and then summed along the diagonals. This summation is done in layers known as reduction layers, and for two numbers of length  $N$ , there are  $\log(N)$  reduction layers, which is good for speeding up the computation.

Direct implementation of lattice multiplication / wallace tree algorithm in a bitwise fashion is highly straightforward:

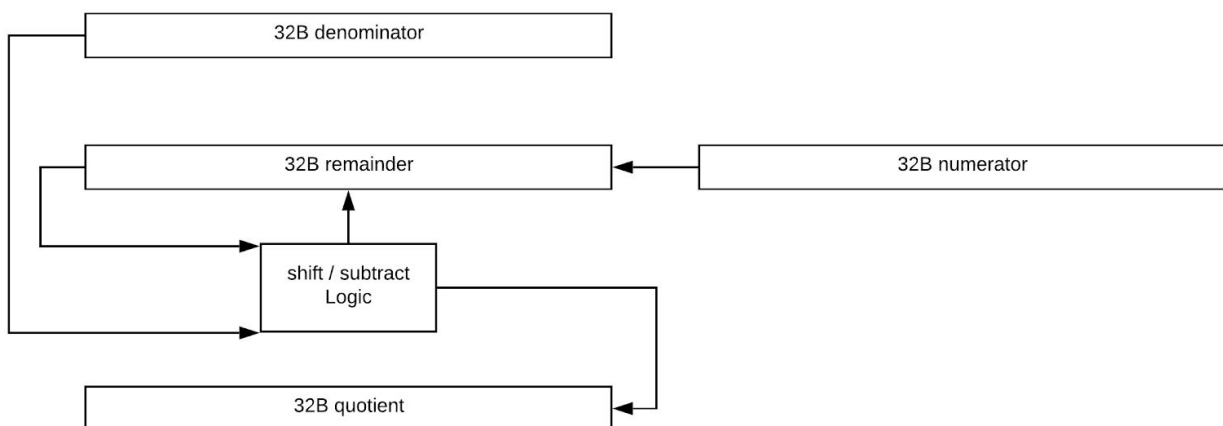
1. Multiply (AND operation) the bits and get a grid or  $N^2$  results.
2. Sum these results diagonally with  $\log(N)$  reduction layers
3. Sum the previous summation results together to get the final multiplication product

However, for 2 32-bit numbers the hardware can grow extremely complex and large, and this leads to the modification of the Wallace tree algorithm which uses BRAM to speed up the computations, using pre-computed results. If instead of performing the lattice operation bit-wise we can perform the operations on 4-bits at a time and use BRAM to store the pre-computed results of such 4-bit operations. This would change the original 5 reduction layers for 32B numbers to 3 reduction layers, leading to shorter critical paths.

The multiplication between each digit could be handled by a multiplication table (in addition to a multiplication “overflow” table) stored in memory of size 16 x 16 to quickly generate the lattice, and summing the results from these digit-wise multiplications over 3 layers should result in paths of lengths that should be able to be computed in a single 100MHz cycle. Assuming the critical paths match this prediction, multiplication should be able to be performed in a single clock cycle.

## Divider

The divider is analogous to a shift-add multiplier but instead of shift-add it is now shift-subtract:



It's algorithm is analogous to the famous long division everyone is taught in grade school. The operation of the divider is best described in pseudo-code:

- If denominator > numerator:
  - Quotient = 0
  - Remainder = numerator
- If denominator == numerator:
  - Quotient = 1
  - Remainder = 0
- Otherwise, repeat 32 times (for 32B number):
  - Quotient[31:1] = quotient[30:0] (shift left)
  - If denominator >= remainder:
    - Remainder[31:1] = remainder - denominator

- Quotient[0] = 1
  - Else:
    - Remainder[31:1] = remainder[30:0] (shift left)
    - Quotient[0] = 0
  - Remainder[0] = numerator[31]
- What remains in the quotient register is the quotient, and what remains in the remainder register is the remainder

The control for the divider is straightforward if none of the initial two “easy” conditions are met, then it’s a simple loop-shift-subtract 32 times. With this setup, the guaranteed worst case performance of the division algorithm is 32 cycles of the clock.